

assignment-1

August 15, 2023

1 1. Importing Libraries

```
[1]: import pandas as pd
import numpy as np
import statistics
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data=pd.read_excel("daily_offers.xlsx",parse_dates=['delivery date'])
```

```
[3]: df=data.copy()
```

2 2. Exploring the dataset

```
[4]: df.head()
```

```
[4]:
```

	id	item_date	quantity	tons	customer	\
0	EC06F063-9DF0-440C-8764-0B0C05A4F6AE	20210401.0	54.151139	30156308.0		
1	4E5F4B3D-DDDF-499D-AFDE-A3227EC49425	20210401.0	768.024839	30202938.0		
2	E140FF1B-2407-4C02-A0DD-780A093B1158	20210401.0	386.127949	30153963.0		
3	F8D507A0-9C62-4EFE-831E-33E1DA53BB50	20210401.0	202.411065	30349574.0		
4	4E1C4E78-152B-430A-8094-ADD889C9D0AD	20210401.0	785.526262	30211560.0		

	country	status	item type	application	thickness	width \
0	28.0	Won	W	10.0	2.00	1500.0
1	25.0	Won	W	41.0	0.80	1210.0
2	30.0	Won	WI	28.0	0.38	952.0
3	32.0	Won	S	59.0	2.30	1317.0
4	28.0	Won	W	10.0	4.00	2000.0

[illegible]

```

    selling_price
0         854.00
1        1047.00
2         644.33
3         768.00
4         577.00

```

```
[5]: df.shape
```

```
[5]: (181673, 14)
```

```
[6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 181673 entries, 0 to 181672
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    181671 non-null  object
1   item_date             181672 non-null  float64
2   quantity tons        181673 non-null  object
3   customer              181672 non-null  float64
4   country               181645 non-null  float64
5   status               181671 non-null  object
6   item type            181673 non-null  object
7   application          181649 non-null  float64
8   thickness            181672 non-null  float64
9   width                181673 non-null  float64
10  material_ref         103754 non-null  object
11  product_ref          181673 non-null  int64
12  delivery date        181672 non-null  object
13  selling_price        181672 non-null  float64
dtypes: float64(7), int64(1), object(6)
memory usage: 19.4+ MB

```

```
[7]: df.corr()
```

```

[7]:
    item_date  customer  country  application  thickness  \
item_date    1.000000 -0.008821 -0.015801   -0.015962   0.003075
customer    -0.008821  1.000000  0.083560    0.000882   0.009589
country     -0.015801  0.083560  1.000000   -0.019350  -0.019580
application -0.015962  0.000882 -0.019350    1.000000  -0.059472
thickness    0.003075  0.009589 -0.019580   -0.059472   1.000000
width        0.020480  0.009203  0.055295   -0.204430   0.161714
product_ref  0.037441 -0.007716 -0.147389   -0.131843   0.038082
selling_price 0.004467 -0.000053  0.002993    0.001462  -0.001130

```

	width	product_ref	selling_price
item_date	0.020480	0.037441	0.004467
customer	0.009203	-0.007716	-0.000053
country	0.055295	-0.147389	0.002993
application	-0.204430	-0.131843	0.001462
thickness	0.161714	0.038082	-0.001130
width	1.000000	-0.034460	0.000583
product_ref	-0.034460	1.000000	0.002118
selling_price	0.000583	0.002118	1.000000

```
[8]: df.describe()
```

```
[8]:
```

	item_date	customer	country	application \
count	1.816720e+05	1.816720e+05	181645.000000	181649.000000
mean	2.020459e+07	3.051221e+07	44.893022	25.615809
std	4.551119e+03	2.433382e+07	24.404214	17.754175
min	1.995000e+07	1.245800e+04	25.000000	2.000000
25%	2.020093e+07	3.019688e+07	26.000000	10.000000
50%	2.020113e+07	3.020524e+07	30.000000	15.000000
75%	2.021020e+07	3.028042e+07	78.000000	41.000000
max	2.021040e+07	2.147484e+09	113.000000	99.000000

	thickness	width	product_ref	selling_price
count	181672.000000	181673.000000	1.816730e+05	1.816720e+05
mean	2.564827	1295.286724	4.739679e+08	1.918036e+03
std	6.572321	261.631754	7.175101e+08	3.317956e+05
min	0.180000	1.000000	6.117280e+05	-1.160000e+03
25%	0.700000	1180.000000	6.119930e+05	6.690000e+02
50%	1.500000	1250.000000	6.406650e+05	8.120000e+02
75%	3.000000	1500.000000	1.332077e+09	9.530000e+02
max	2500.000000	2990.000000	1.722208e+09	1.000010e+08

3 3. Data Cleaning

3.1 3.1 Dropping Irrelevant Columns

```
[9]: df=df.drop(columns=['id','item_date','material_ref','product_ref'],axis=1)
```

```
[10]: df.head(1)
```

```
[10]:
```

	quantity tons	customer	country	status	item type	application	thickness \
0	54.151139	30156308.0	28.0	Won	W	10.0	2.0

	width	delivery date	selling_price
0	1500.0	20210701	854.0

3.2 3.2 Handling Missing Values

```
[11]: df.isnull().sum()
```

```
[11]: quantity tons      0
      customer         1
      country         28
      status          2
      item type        0
      application      24
      thickness        1
      width           0
      delivery date    1
      selling_price    1
      dtype: int64
```

```
[12]: for x in df.columns:
      null=df[x].isnull().sum()
      percentage=(null/len(df))*100
      print(x, ' : ', "%.3f"%percentage, '%')
```

```
quantity tons : 0.000 %
customer : 0.001 %
country : 0.015 %
status : 0.001 %
item type : 0.000 %
application : 0.013 %
thickness : 0.001 %
width : 0.000 %
delivery date : 0.001 %
selling_price : 0.001 %
```

4 Filling missing values

To fill the nan values we find the which entity have maximum frequency in each column to find frequency we calculate the mode

```
[13]: df['country'].fillna(df['country'].mode()[0], inplace=True)
      df['application'].fillna(df.application.mode()[0], inplace=True)
      df['customer'].fillna(df['customer'].mode()[0], inplace=True)
      df['thickness'].fillna(df['thickness'].mean(), inplace=True)
      df['selling_price'].fillna(df['selling_price'].mean(), inplace=True)
      df['delivery date'].fillna(df['delivery date'].mode()[0], inplace=True)
      df['status'].fillna(df['status'].mode()[0], inplace=True)
```

```
[14]: df.isna().sum()
```

```
[14]: quantity tons      0
      customer          0
      country           0
      status            0
      item type         0
      application       0
      thickness         0
      width             0
      delivery date     0
      selling_price     0
      dtype: int64
```

```
[15]: df.dropna(inplace=True)
      df.isna().sum()
```


```
[15]: quantity tons      0
      customer          0
      country           0
      status            0
      item type         0
      application       0
      thickness         0
      width             0
      delivery date     0
      selling_price     0
      dtype: int64
```

```
[16]: df.drop(df[df['quantity tons']=='e'].index,inplace=True) # dataset had a record
      ↪with id='e'
```

4.1 3.3 Handling Outliers

```
[17]: plt.figure(figsize = (10,5))
      df.boxplot(column=['thickness','width','selling_price'])
```

```
[17]: <AxesSubplot:>
```



```
[18]: df['zscore']=(df['selling_price']-df['selling_price'].mean())/
      ↪df['selling_price'].std()
      df
```

```
[18]:
```

	quantity tons	customer	country	status	item type	application \
0	54.151139	30156308.0	28.0	Won	W	10.0
1	768.024839	30202938.0	25.0	Won	W	41.0
2	386.127949	30153963.0	30.0	Won	WI	28.0
3	202.411065	30349574.0	32.0	Won	S	59.0
4	785.526262	30211560.0	28.0	Won	W	10.0
...
181668	102.482422	30200854.0	25.0	Won	W	41.0
181669	208.086469	30200854.0	25.0	Won	W	41.0
181670	4.235594	30200854.0	25.0	Won	W	41.0
181671	-2000	30200854.0	25.0	Won	W	41.0
181672	406.686538	30200854.0	25.0	Won	W	41.0

	thickness	width	delivery date	selling_price	zscore
0	2.00	1500.0	20210701	854.00	-0.003207
1	0.80	1210.0	20210401	1047.00	-0.002625
2	0.38	952.0	20210101	644.33	-0.003839
3	2.30	1317.0	20210101	768.00	-0.003466
4	4.00	2000.0	20210301	577.00	-0.004042
...
181668	0.96	1220.0	20200701	591.00	-0.004000
181669	0.95	1500.0	20200701	589.00	-0.004006
181670	0.71	1250.0	20200701	619.00	-0.003915


181671	0.85	1250.0	20200701	601.00	-0.003969
181672	0.71	1240.0	20200701	607.00	-0.003951

[181672 rows x 11 columns]

```
[19]: df=(df[(df['zscore']>=-3) & (df['zscore']<=3)])
```


```
[20]: df.boxplot(column=['thickness','width','selling_price'])
```

[20]: <AxesSubplot:>




```
[21]: c=df.value_counts(df['status'])
def autopct(pct):
    return('%.1f'%pct) if pct>1 else ''
label=['Won','Lost','Not lost for AM ','Revised','To be_
approved','Draft','Offered','Offerable','Wonderful']
plt.figure(figsize = (9,9))
plt.pie(c,autopct=autopct)
plt.legend(label,loc='best',bbox_to_anchor=(1.5, 0.8))
```

[21]: <matplotlib.legend.Legend at 0x282d47c5220>



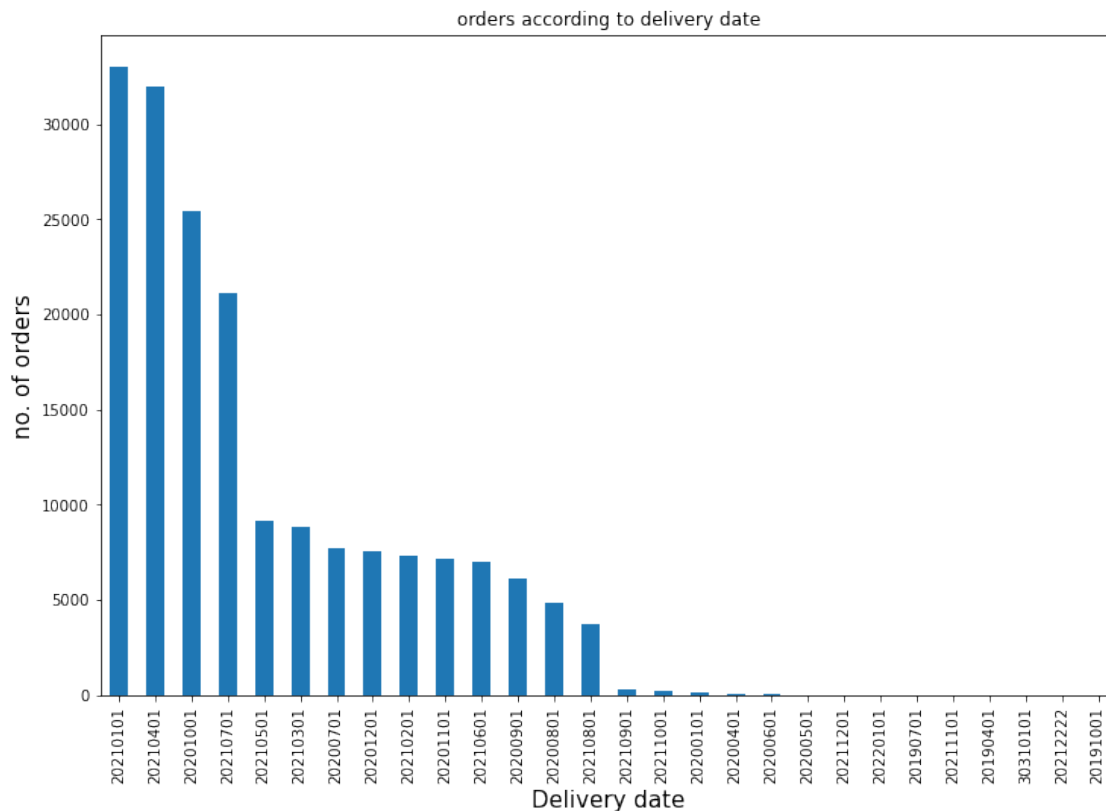
```
[22]: x=df['selling_price'].groupby(df['item type']).value_counts().unstack('item_
      ↪type').sum(axis=0)
      plt.figure(figsize = (15,5))
      plt.plot(x,marker='.',color='purple')
      plt.xlabel('item type',fontsize=15)
      plt.ylabel('selling price',fontsize=15)
```

[22]: Text(0, 0.5, 'selling price')





```
[23]: x=df['delivery date'].value_counts()
ax=x.plot(kind='bar',ylabel='orders',title='orders according to delivery_
↳date',figsize = (12,8))
ax.set_xlabel('Delivery date',fontsize=15)
ax.set_ylabel('no. of orders',fontsize=15)
```

```
[23]: Text(0, 0.5, 'no. of orders')
```



```
[24]: k=df['item type'].value_counts()
def autopct(pct):
    return ('%.2f' %pct) if pct > 1 else ''
plt.figure(figsize = (9,9))
plt.pie(k,explode=[0.1,0,0,0.3,0.3,0,0],autopct=autopct,colors=sns.
↳color_palette('flare'))
plt.legend(k.index)
```

```
[24]: <matplotlib.legend.Legend at 0x282d6300670>
```



```
[25]: df_model=df[['quantity tons','country','status','item_
↳type','application','thickness','width','selling_price']]
```

```
[26]: for x in df_model.columns:
        print(x," has ",len(df[x].unique())," categories")
```

```
quantity tons has 181670 categories
country has 17 categories
status has 9 categories
item type has 7 categories
application has 30 categories
thickness has 595 categories
width has 1386 categories
selling_price has 9794 categories
```

```
[27]: top10=[x for x in df_model['country'].value_counts().
        ↪sort_values(ascending=False).head(15).index]
for x in top10:
    df_model[x]=np.where(df_model['country']==x,1,0)
```

C:\Users\pulki\AppData\Local\Temp\ipykernel_10032\3970385872.py:3:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_model[x]=np.where(df_model['country']==x,1,0)
C:\Users\pulki\AppData\Local\Temp\ipykernel_10032\3970385872.py:3:  

    SettingWithCopyWarning:  

    A value is trying to be set on a copy of a slice from a DataFrame.  

    Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_model[x]=np.where(df_model['country']==x,1,0)
```

```
[28]: top10=[x for x in df_model['status'].value_counts().
        ↪sort_values(ascending=False).head(6).index]
def onehot(df,variable,top_labels):

    for label in top_labels:
        df_model[label]=np.where(df_model[variable]==label,1,0)

onehot(df_model,'status',top10)
```

```
[29]: top10=[x for x in df_model['item type'].value_counts().
        ↪sort_values(ascending=False).head(7).index]
for x in top10:
    df_model[x]=np.where(df_model['item type']==x,1,0)
```

```
[30]: top10=[x for x in df_model['application'].value_counts().
        ↪sort_values(ascending=False).head(15).index]
for x in top10:
    df_model[x]=np.where(df_model['application']==x,1,0)
```

```
[31]: df_model.drop(['country','status','item_
        ↪type','application'],axis=1,inplace=True)
```

```
[32]: plt.figure(figsize=(25,15))
sns.heatmap(df_model.corr(),annot=True,cmap="RdYlGn",cbar=True)
```

thickness	116	0.16	0.0970	0.40	0.0780	0.2040	0.08	0.0770	0.0280	0.230	0.0430	0.160	0.2000	0.038	0.026	0.03	0.170	0.0420	0.060	0.020	0.0980	0.180	0.160	0.2040	0.0220	0.10	0.024	0.18	0.00	0.0080	0.07	0.12	0.14	0.0330	0.0380	0.0280	0.0980	0.0240	0.14		
width	116	0.1	0.18	0.0010	0.3550	0.0890	0.0610	0.30	0.0630	0.02	0.0650	0.0490	0.02	0.0780	0.14	0.0630	0.0630	0.08	0.12	0.12	0.0358	0.11	0.0920	0.08	0.0630	0.070	0.20	0.0490	0.18	0.0044	12	0.15	0.22	0.0170	0.0820	0.170	0.120	0.080	0.52		
selling price	0.9794	12	0.1	0.03	0.06	0.0280	0.0440	0.150	0.0290	0.0400	0.0420	0.230	0.0210	0.170	0.0120	0.0680	0.0640	0.170	0.034	0.02	0.0558	0.36	0.065	0.40	0.0620	0.130	0.0280	0.030	0.081	0.12	0.1	0.760	0.080	0.0560	0.0780	0.20	0.13				
75	0.4	0.04	0.0014	0.3	0.1	0.0280	0.0280	0.36	0.16	0.2	0.13	0.1	0.4	0.670	0.040	0.450	0.04	0.039	0.02	0.13	0.1	0.10	0.650	0.060	0.1020	0.59	0.13	0.16	0.60	0.0000	0.0510	0.034	0.18	0.12	0.4	0.10	0.040	0.170	0.40	0.03	0.075
26	0.6	0.078	0.30	0.060	0.2	0.1	0.0087	0.0980	0.020	0.0340	0.0210	0.10	0.020	0.290	0.02	0.060	0.290	0.02	0.0770	0.10	0.10	0.078	0.11	0.10	0.058	0.058	0.0020	0.080	0.040	0.0240	0.10	0.16	0.0980	0.060	0.0780	0.20	0.13				
25	0.4	0.240	0.0850	0.10	0.028	0.08	0.1	0.0993	0.108	0.088	0.058	0.10	0.24	0.10	0.088	0.098	0.064	0.058	0.070	0.2	0.22	0.048	0.109	0.048	0.04	0.0698	0.10	0.0530	0.37	0.0100	0.0200	0.0790	0.530	0.330	0.210	0.170	0.108	0.068	0.055		
27	0.4	0.040	0.068	0.042	0.40	0.098	0.059	0.1	0.120	0.098	0.270	0.180	0.0280	0.108	0.088	0.098	0.098	0.064	0.068	0.080	0.2	0.028	0.0570	0.0540	0.10	0.0440	0.120	0.180	0.120	0.0880	0.060	0.580	0.030	0.210	0.180	0.160	0.10	0.068	0.065		
32	0.70	0.50	0.030	0.10	0.160	0.038	0.18	0.12	0.1	0.260	0.090	0.080	0.760	0.490	0.30	0.30	0.0930	0.020	0.21	0.1	0.1	0.170	0.0820	0.430	0.080	0.250	0.04	0.07	0.1930	0.060	0.04	0.0180	0.160	0.120	0.320	0.30	0.098	0.042	0.0920	0.250	0.12
28	0.4	0.230	0.630	0.290	0.20	0.088	0.098	0.098	0.028	0.1	0.260	0.290	0.04	0.108	0.098	0.098	0.058	0.180	0.240	0.050	0.380	0.098	0.120	0.090	0.078	0.105	0.058	0.088	0.0020	0.080	0.550	0.340	0.210	0.170	0.108	0.098	0.068	0.061			
84	0.8	0.0880	0.2	0.048	0.130	0.260	0.250	0.20	0.930	0.02	0.1	0.0650	0.62	0.40	0.31	0.20	0.260	0.0190	0.170	0.340	0.12	0.4	0.0630	0.370	0.0390	0.790	0.050	0.040	0.108	0.008	0.008	0.065	0.1	0.0340	0.430	0.25	0.03	0.0880	0.10	0.075	
39	0.4	0.040	0.260	0.242	0.110	0.0440	0.180	0.040	0.180	0.080	0.290	0.06	0.1	0.0530	0.040	0.220																									

```
# Independent and dependent features
X=df_model.drop("selling_price",axis=1)
Y=df_model.selling_price
```

```
# train test split
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.
↳1,random_state=10)
```

```
## standardising dataset
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

```
X_train=scaler.fit_transform(X_train)
```

```
C:\Users\pulki\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688:
FutureWarning: Feature names only support names that are all strings. Got
feature names with dtypes: ['float', 'str']. An error will be raised in 1.2.
  warnings.warn(
C:\Users\pulki\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688:
FutureWarning: Feature names only support names that are all strings. Got
```

```
feature names with dtypes: ['float', 'str']. An error will be raised in 1.2.  
warnings.warn(  

```

```
[37]: X_test=scaler.transform(X_test)
```

```
C:\Users\pulki\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688:  
FutureWarning: Feature names only support names that are all strings. Got  
feature names with dtypes: ['float', 'str']. An error will be raised in 1.2.  
warnings.warn(  

```

4.1.1 Model 1: Linear Regression

```
[38]: from sklearn.linear_model import LinearRegression  
      regression=LinearRegression()
```

```
[39]: regression.fit(X_train,Y_train)
```

```
[39]: LinearRegression()
```


```
[40]: reg_pred=regression.predict(X_test)
```

```
[41]: print("coefficient: ",regression.coef_)
```

```
coefficient: [-2.32672102e-01 -1.62186773e+01 -2.53134941e+01 -1.62463257e+01  
-7.74571319e-02 -3.75626623e+00  5.43705972e-01  3.59723804e+00  
-7.73219377e+00  1.25464816e+01 -1.28449933e+00  5.69921797e-01  
 1.27372860e+00  2.91647142e+00  4.06125890e+00 -5.07593388e+00  
-3.29474824e+00  1.47935036e+01 -7.43696323e+01 -5.95753202e+01  
-4.15045321e+01 -2.04294660e+01 -1.44457259e+01 -1.27583167e+02  
-5.88441147e+13 -5.79300202e+13 -2.07318321e+13 -6.89344364e+12  
-6.35159796e+12 -1.41451754e+12 -2.94967141e+11  5.40022430e+00  
 6.34479696e+01 -6.04193636e+00 -1.53455089e+01  1.87666388e+01  
-1.30293093e+01  3.82073744e+00  7.18635847e-01  3.01577521e+00]
```

```
[42]: plt.scatter(Y_test,reg_pred)
```

```
[42]: <matplotlib.collections.PathCollection at 0x282d71762e0>
```



```
[43]: from sklearn.metrics import mean_squared_error
      print(mean_squared_error(Y_test,reg_pred))
```

32212.770174482248

4.1.2 Model 2: Random Forest Regression

```
[44]: from sklearn.ensemble import RandomForestRegressor
      rf=RandomForestRegressor(n_estimators=100,random_state=42)
```

```
[ ]: rf.fit(X_train,Y_train)
```

```
[ ]: rf_pred=rf.predict(X_test)
```

```
[ ]: print(mean_squared_error(Y_test,rf_pred))
```

```
[ ]: plt.scatter(Y_test,rf_pred)
```

4.1.3 Model 3: Lasso Regression

```
[ ]: from sklearn.linear_model import Lasso
      lm=Lasso(alpha=0.1)
      lm.fit(X_train,Y_train)
      lm_pred=lm.predict(X_test)
```

```
[ ]: print(mean_squared_error(Y_test,lm_pred))
```

```
[ ]: plt.scatter(Y_test,lm_pred)
```

4.1.4 Model 4: Ridge Regression

```
[ ]: from sklearn.linear_model import Ridge
rm=Ridge(alpha=0.1)
rm.fit(X_train,Y_train)
```

```
[ ]: rm_pred=rm.predict(X_test)
```

```
[ ]: print(mean_squared_error(Y_test,lm_pred))
```

```
[ ]: plt.scatter(Y_test,rm_pred)
```

```
[ ]: from sklearn.ensemble import RandomForestRegressor
rf_random=RandomForestRegressor()
```

```
[ ]: from sklearn.metrics import r2_score
score=r2_score(Y_test, rf_pred)
```

```
[ ]: predictions=rf_random.predict(X_test)
```

```
[ ]: #Hyperparameter Tuning Randomforest

# Randomised search CV
#Number of trees in random forest
n_estimators=[int(x) for x in np.linspace(start=100,stop=1200,num=12)]
#no. of features to consider at every split
max_features=['auto','sqrt']
#max no. of levels in tree
max_depth=[int(x) for x in np.linspace(5,30,num=6)]
# Min no. of samples required to split a node
min_samples_split=[2,5,10,15,100]
min_samples_leaf=[1,2,5,10]
```

```
[ ]: from sklearn.model_selection import RandomizedSearchCV
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf}

print(random_grid)
```

```
[ ]: rf_random = RandomizedSearchCV(estimator = rf, param_distributions =  
    ↪random_grid,scoring='neg_mean_squared_error', n_iter = 10, cv = 5,  
    ↪verbose=2, random_state=42,n_jobs=1)
```

```
[ ]: plt.scatter(Y_test,predictions)
```

```
[ ]: sns.distplot(Y_test-predictions)
```

```
[ ]: rf_random.fit(X_train,Y_train)
```